

Docker Image Security Best Practices

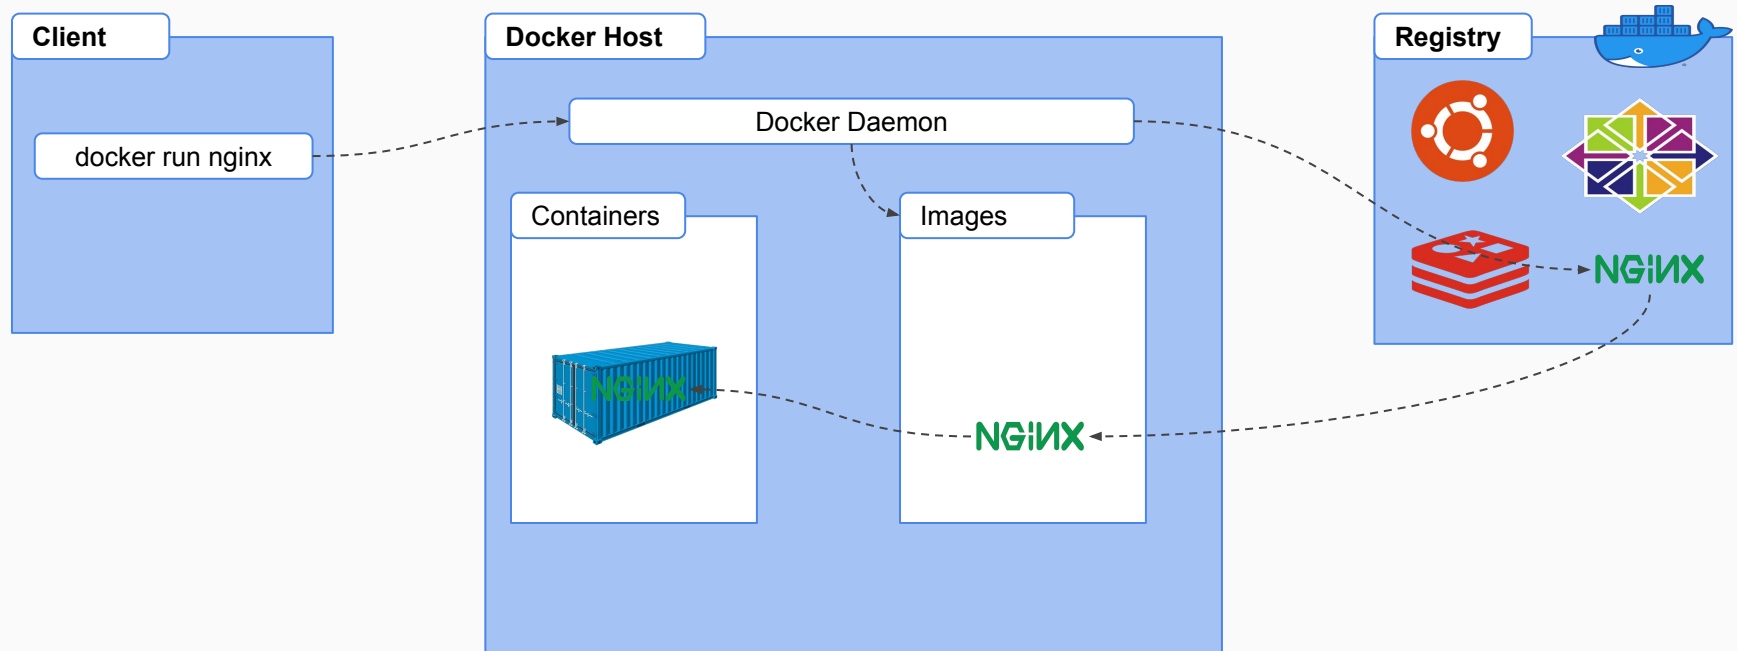
By Akshay Ithape - DevOps Engineer

\$whoami

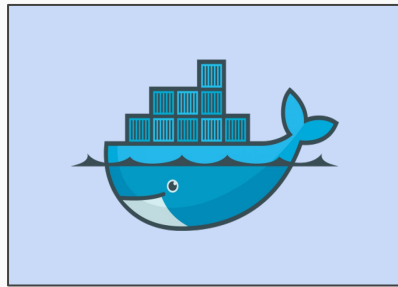


- Myself Akshay Ithape.
- Currently, I'm Working as a DevOps Engineer At Eastern Enterprise.
- I'm a passionate about Linux and Cloud.
- As a Linux Person, I believes in Open Source so I like to share my knowledge with community.
- I'm RedHat, AWS, Terraform and Kubernetes Certified Engineer.
- Written articles For "Open Source For You" Magazine.

How Docker Works ?



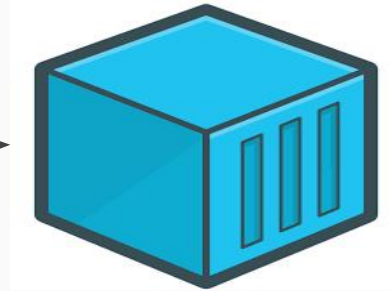
Why Docker Image is Required ?



Docker Image

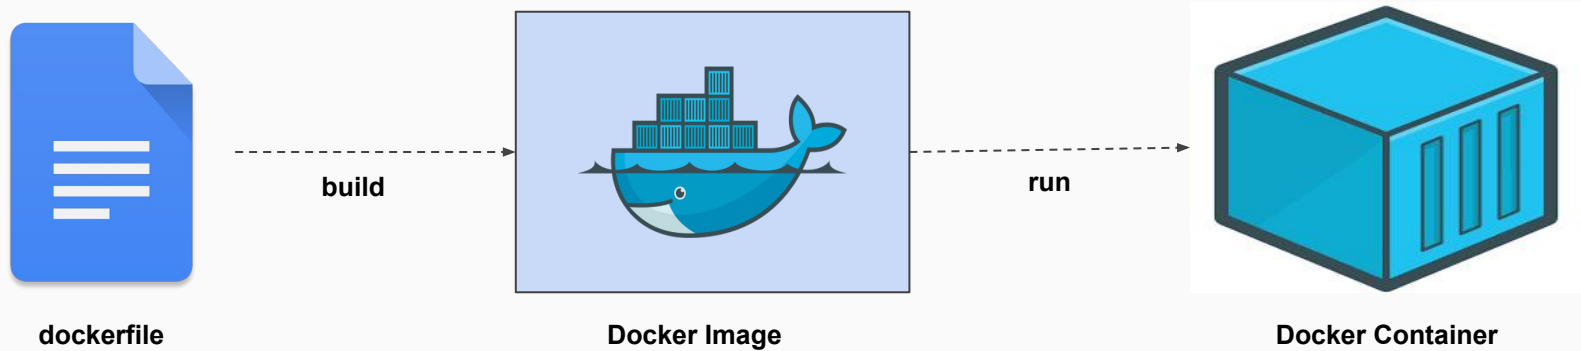


run



Docker Container

How To Build Custom Docker Image ?



Simple Dockerfile Context

```
dockerfile-format x
dockerfile-format
1  # Comment
2  INSTRUCTION arguments
```

```
dockerfile x
dockerfile
1  FROM nginx:latest
2  RUN apt-get update -y
3  WORKDIR /usr/share/nginx/html
4  COPY index.html .
5  CMD ["nginx", "-g", "daemon off;"]
6  EXPOSE 80
```

Choose Official or Verified Base Image

Always choose Official or Verified base image. If is not official or verified then vulnerability chances are very high.

The screenshot displays the Docker Hub search results for the 'nginx' image. The top section shows three specific image entries:

- cloudctl/nginx**: By cloudctl • Updated 2 hours ago. 2.1K Downloads, 0 Stars. Tags: Container, Linux, x86-64.
- nginx** (Official Image): Updated 6 hours ago. 1B+ Downloads, 10K+ Stars. Official build of Nginx. Tags: Container, Linux, 386, PowerPC 64 LE, IBM Z, mips64le, x86-64, ARM, ARM 64, Application Infrastructure.
- bitnami/nginx** (Verified Publisher): By Bitnami • Updated 5 hours ago. 100M+ Downloads, 120 Stars. Bitnami nginx Docker Image. Tags: Container, Linux, x86-64.

The bottom section shows the Docker Hub search interface with the search bar containing 'nginx' and the 'Containers' tab selected. The search results are filtered to show 'Verified Publisher' and 'Official Images' only. The first result is the 'nginx' official image, updated 6 hours ago.

Prefer minimal base images

- Choose images with fewer OS libraries and tools lower the risk and attack surface of the containers.
- Prefer alpine-based images over full-blown system OS images.

TAG			docker pull nginx:latest
latest  Log4Shell CVE not detected			
Last pushed a day ago by doijanky			
DIGEST	OS/ARCH	COMPRESSED SIZE	
18100eaa5648	linux/386	55.94 MB	
1a763cbd30ef	linux/amd64	54.1 MB	
bcba09dc6c63	linux/arm/v5	50.97 MB	

TAG			docker pull nginx:alpine
alpine  Log4Shell CVE not detected			
Last pushed a day ago by doijanky			
DIGEST	OS/ARCH	COMPRESSED SIZE	
2cb22820489f	linux/386	10.13 MB	
1e3458b88413	linux/amd64	9.69 MB	
efddee619e0f	linux/arm/v6	9.22 MB	

Vulnerabilities Scan Reports

```
Package manager:  deb
Project name:     docker-image|nginx
Docker image:     nginx:latest
Platform:        linux/amd64
Base image:       nginx:1.21.6
```

Tested 143 dependencies for known vulnerabilities, **found 84 vulnerabilities.**

According to our scan, you are currently using the most secure version of the selected base image

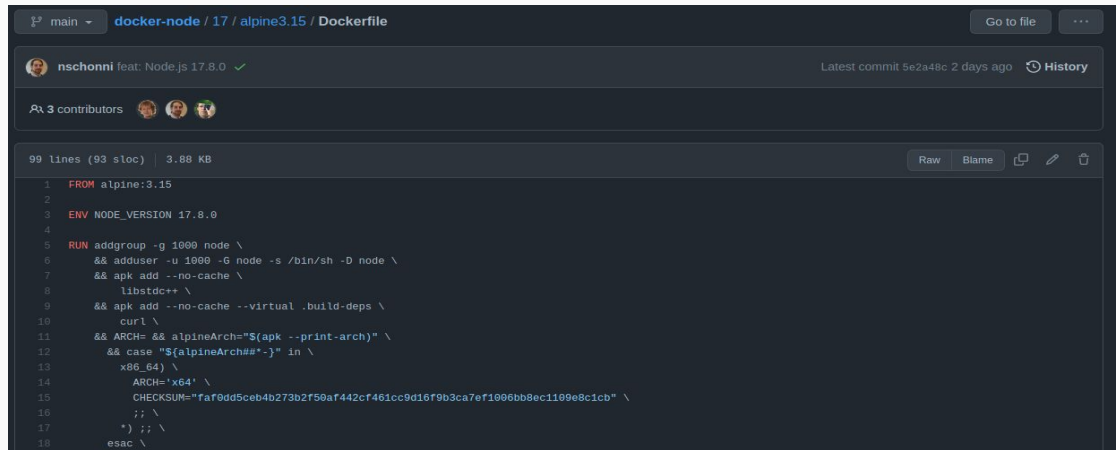
For more free scans that keep your images secure, sign up to Snyk at <https://dockr.ly/3ePqVcp>

```
Package manager:  apk
Project name:     docker-image|nginx
Docker image:     nginx:alpine
Platform:        linux/amd64
Base image:       nginx:1.21.6-alpine
```

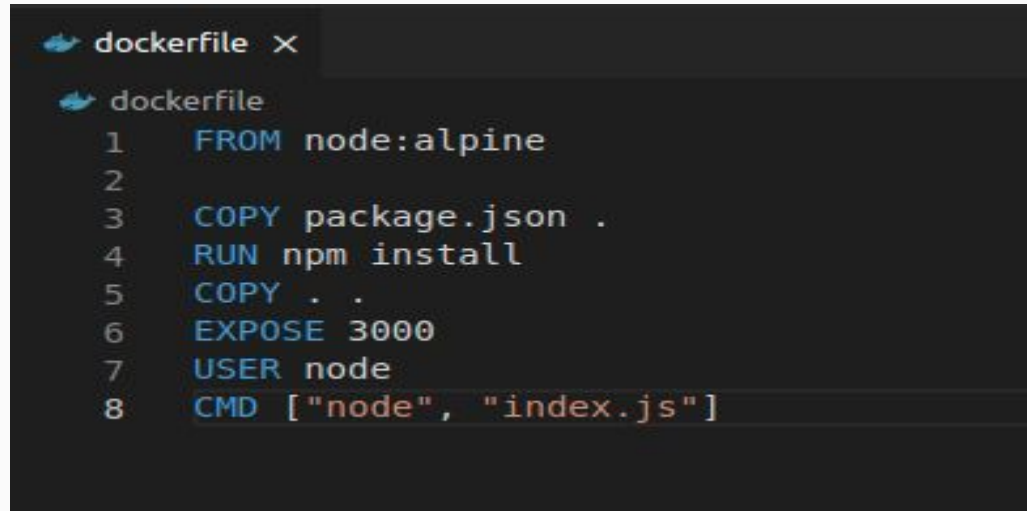
Tested 43 dependencies for known vulnerabilities, **found 1 vulnerability.**

Least privileged user

Create a dedicated user and group on the image, with minimal permissions to run the application; use the same user to run this process.

A screenshot of a Dockerfile repository on GitHub. The repository is named 'docker-node' and is located at path '/ 17 / alpine3.15 / Dockerfile'. The commit is by 'nschonni' with the message 'feat: Node.js 17.8.0'. The Dockerfile content is as follows:

```
1 FROM alpine:3.15
2
3 ENV NODE_VERSION 17.8.0
4
5 RUN addgroup -g 1000 node \
6     && adduser -u 1000 -G node -s /bin/sh -D node \
7     && apk add --no-cache \
8         libstdc++ \
9         && apk add --no-cache --virtual .build-deps \
10             curl \
11             && ARCH= && alpineArch="$(apk --print-arch)" \
12             && case "${alpineArch}" in \
13                 x86_64) \
14                     ARCH='x64' \
15                     CHECKSUM='fa0dd5ceb4b273b2f50af442cf461cc9d16f9b3ca7ef1066bb8ec1109e8c1cb' \
16                     ;; \
17                 *) ;; \
18             esac \
```

A screenshot of a Dockerfile in a code editor. The file is named 'dockerfile'. The content is as follows:

```
1 FROM node:alpine
2
3 COPY package.json .
4 RUN npm install
5 COPY . .
6 EXPOSE 3000
7 USER node
8 CMD ["node", "index.js"]
```

<https://www.linkedin.com/in/akshayithape/>

Verify and Sign images to mitigate MITM attacks

We put a lot of trust into docker images. It is critical to make sure the image we're pulling is the one pushed by the publisher, and that no one has tampered with it.

```
> export DOCKER_CONTENT_TRUST=1
```

```
> docker pull vigneshkumar73/vicky_nginx
```

```
Using default tag: latest
```

```
Error: remote trust data does not exist for docker.io/vigneshkumar73/vicky_nginx: notary.docker.io does not have trust data for docker.io/vigneshkumar73/vicky_nginx
```

```
> docker trust key generate dev1
```

```
Generating key for dev1...
```

```
Enter passphrase for new dev1 key with ID 1ebd66c:
```

```
Repeat passphrase for new dev1 key with ID 1ebd66c:
```

```
Successfully generated and loaded private key. Corresponding public key available: /home/akshay/docker-test/dev1.pub
```

```
> docker tag node-test:v1 imperishableakki/node-test:v1
```

```
> docker trust signer add --key dev1.pub dev1 imperishableakki/node-test
```

```
Adding signer "dev1" to imperishableakki/node-test...
```

```
Initializing signed repository for imperishableakki/node-test...
```

```
You are about to create a new root signing key passphrase. This passphrase will be used to protect the most sensitive key in your signing system. Please choose a long, complex passphrase and be careful to keep the password and the key file itself secure and backed up. It is highly recommended that you use a password manager to generate the passphrase and keep it safe. There will be no way to recover this key. You can find the key in your config directory.
```

```
Enter passphrase for new root key with ID 18b7bea:
```

```
Repeat passphrase for new root key with ID 18b7bea:
```

```
Enter passphrase for new repository key with ID aeaa583:
```

```
Repeat passphrase for new repository key with ID aeaa583:
```

```
Successfully initialized "imperishableakki/node-test"
```

```
Successfully added signer: dev1 to imperishableakki/node-test
```

```
> docker trust sign imperishableakki/node-test:v1
```

```
Signing and pushing trust data for local image imperishableakki/node-test:v1, may overwrite remote trust data
```

```
The push refers to repository [docker.io/imperishableakki/node-test]
```

```
45f534bae6c2: Pushed
```

```
e5623c90b52f: Pushed
```

```
80f4d40e1c68: Pushed
```

```
0ad6919e1cc3: Mounted from library/node
```

```
5a09a182660a: Mounted from library/node
```

```
41c27a423d25: Mounted from library/node
```

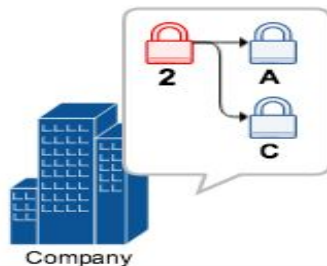
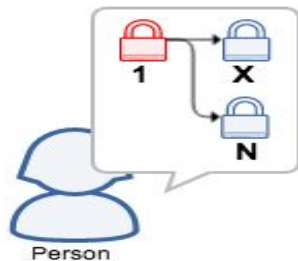
```
ff768a1413ba: Mounted from library/node
```

```
v1: digest: sha256:d7e67a0b2cfab1476adbe69c3db2927f0b644ca56e7593111d9bad8d78a14d2d size: 1780
```

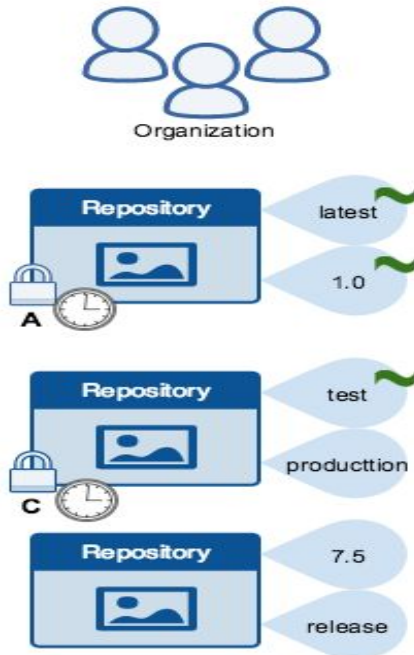
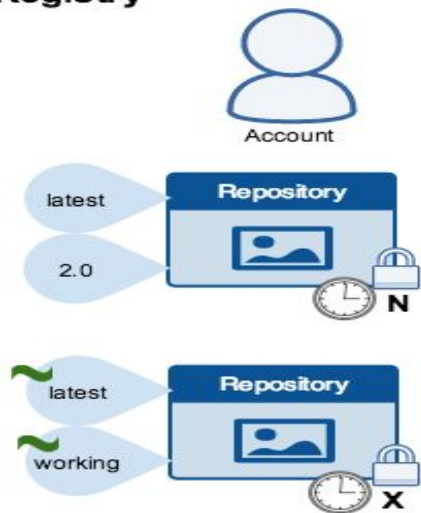
```
Signing and pushing trust metadata
```

```
Enter passphrase for dev1 key with ID 1ebd66c:
```

```
Successfully signed docker.io/imperishableakki/node-test:v1
```



Registry



A offline key is used to create tagging keys. Offline keys belong to a person or an organization. Resides client-side. You should store these in a safe place and back them up.



A tagging key is associated with an image repository. Creators with this key can push or pull any tag in this repository. This resides on client-side.



Timestamp Key

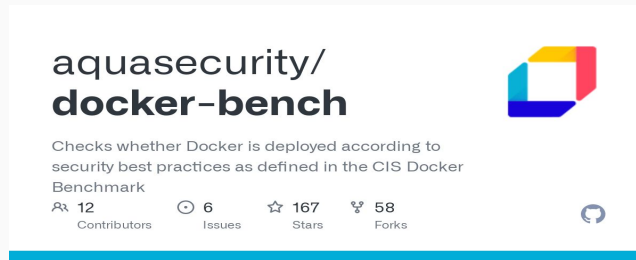
A timestamp key is associated with an image repository. This is created by Docker and resides on the server.



Signed tag.

Find the docker image vulnerabilities

Scan your docker images
for known vulnerabilities
and integrate it as part of
your continuous
integration. There are
many open sources
available to scan images.



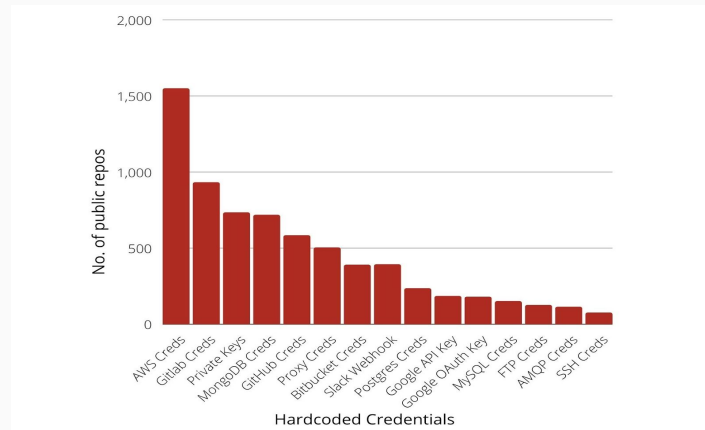
& Many More!

Trivy Demo



Don't leak sensitive information to docker images

It's easy to accidentally leak secrets, tokens, and keys into images when building them.



Top 5 Exposures in Docker Images

- **Hardcoded secrets**
- **Sensitive config files**
- **Adding the entire git repo**
- **Paid software licenses**
- **Default credentials**

<https://redhuntlabs.com/blog/scanning-millions-of-publicly-exposed-docker-containers-thousands-of-secrets-leaked.html>

46076

Docker Containers

Leaked at least one
Hardcoded Secret or
Config file

Exposures in Docker
Containers

15,541

Hardcoded Secrets were
identified across 10,181
Repositories

57,589

Potentially Sensitive Config
files copied to Docker
Images across 36,176 Repos

How to proactively stop exposures in docker images?

- ❖ Don't hardcode tokens/API keys in docker images
- ❖ Do not clone/download the required files using credentials. Instead copy them to the image.
- ❖ Used .dockerignore file
- ❖ Multi-Stage build
- ❖ Used container private registry.

FileEditSelectionViewGoRunTerminalHelp

EXPLORER

DOCKER-TEST

.dockerignore

.env

dev1.pub

dockerfile

index.js

package.json

private.key

dockerfile

1FROM node:alpine

2

3RUN mkdir /app

4WORKDIR /app

5COPY package.json .

6RUN npm install

7COPY . .

8EXPOSE 3000

9USER node

10CMD ["node", "index.js"]

.dockerignore

1.env

2private.key

PROBLEMS

OUTPUT

DEBUG CONSOLE

TERMINAL

at 09:10:11 am

~/docker-test

> docker build -f dockerfile -t node-test:v3 .

Sending build context to Docker daemon 6.144kB

Step 1/9 : FROM node:alpine

----> 15c6163d954c

Step 2/9 : RUN mkdir /app

----> Using cache

----> 2366349bdb11

Step 3/9 : WORKDIR /app

----> Using cache

----> 4118da81c329

Step 4/9 : COPY package.json .

----> Using cache

----> 81c8480f487b

Step 5/9 : RUN npm install

----> Using cache

----> b4290cd49163

Step 6/9 : COPY . .

----> Using cache

----> 14e1829510c4

Step 7/9 : EXPOSE 3000

----> Using cache

----> 6887a76f2bda

Step 8/9 : USER node

----> Using cache

----> c4952af36f6e

Step 9/9 : CMD ["node", "index.js"]

----> Using cache

----> d8561ad14f9c

Successfully built d8561ad14f9c

Successfully tagged node-test:v3

~/docker-test

> /setup-run.sh

at 09:10:43 am

~/docker-test

> docker run -it node-test:v3 /bin/sh

/app \$ ls -la

total 32

drwxr-xr-x 1 root root 4096 Mar 25 03:39 .

drwxr-xr-x 1 root root 4096 Mar 25 03:40 ..

-rw-rw-r-- 1 root root 16 Mar 25 03:32 .dockerignore

-rw----- 1 root root 190 Mar 24 17:07 dev1.pub

-rw-rw-r-- 1 root root 137 Mar 25 03:39 dockerfile

-rw-rw-r-- 1 root root 287 Mar 24 13:57 index.js

-rw-r--r-- 1 root root 245 Mar 25 03:39 package-lock.json

-rw-rw-r-- 1 root root 473 Mar 24 13:57 package.json

/app \$

zsh

zsh

docker

OUTLINE

Use fixed tags for immutability

Docker image owners can push new versions to the same tags, which may result in inconsistent images during builds, and makes it hard to track if a vulnerability has been fixed.

The following table summarizes the data shown in the three screenshots:

Tag	Log4Shell CVE	Last pushed	by	Digest	OS/ARCH	Compressed Size
alpine3.15	Not detected	a day ago	dojanky	56d1a2b5d6d0	linux/amd64	48.64 MB
				9ddb9e866d8a	linux/arm/v6	47.56 MB
				bdfd93ea78e9	linux/arm/v7	46.9 MB
				+3 more...		
alpine3.14	Not detected	a day ago	dojanky	5d9f6ece78fb	linux/amd64	48.64 MB
				927e1710e913	linux/arm/v6	47.55 MB
				fb14e42ae506	linux/arm/v7	46.9 MB
				+3 more...		
alpine	Not detected	a day ago	dojanky	56d1a2b5d6d0	linux/amd64	48.64 MB
				9ddb9e866d8a	linux/arm/v6	47.56 MB
				bdfd93ea78e9	linux/arm/v7	46.9 MB
				+3 more...		

```
dockerfile x
dockerfile
1 FROM node@sha256:5d9f6ece78fb5fa95be2875d74b6407f06f56b04ca6ecc517476f79dac8f1aaf
2
3 COPY package.json .
4 RUN npm install
5 COPY . .
6 EXPOSE 3000
7 USER node
8 CMD ["node", "index.js"]
```

Use COPY instead of ADD

Arbitrary URLs specified for ADD could result in MITM attacks, or sources of malicious data. In addition, ADD implicitly unpacks local archives which may not be expected and result in path traversal and Zip Slip vulnerabilities.

ADD instructions is more capable the COPY.

- It can handle remote URLs.
- It can auto-extract tar files.

As we using remote URLs so chances are high for MITM attacks or malicious data.

Space and image layer considerations

When local archives are used, ADD automatically extracts them to the destination directory. While this may be acceptable, it adds the risk of zip bombs and Zip Slip vulnerabilities that could then be triggered automatically.

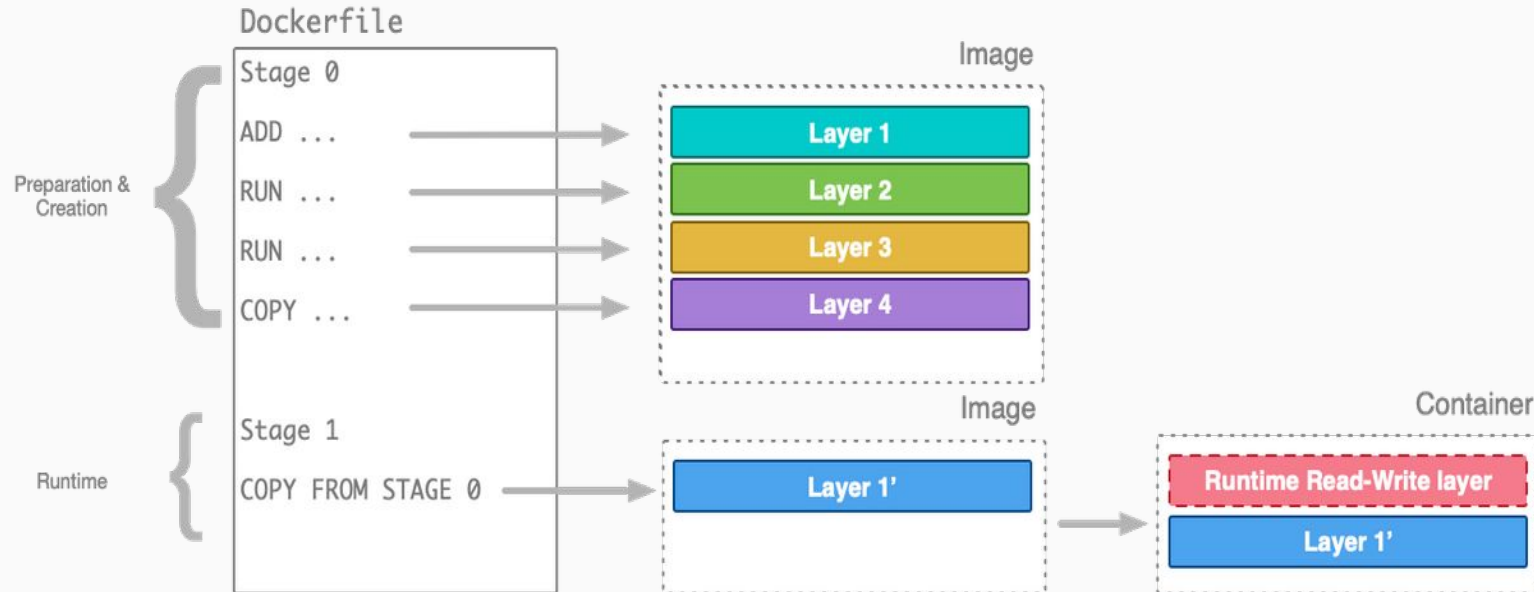
Use labels for metadata

Labels with metadata for images provide useful information for users. Include security details as well.

```
dockerfile x
dockerfile > ...
1 FROM node:alpine3.14
2 LABEL maintainer="Akshay Ithape"
3 LABEL description="Node Apline Image 3.14 for Demo Application"
4
5 RUN mkdir /app
6 WORKDIR /app
7 COPY package.json .
8 RUN npm install
9 COPY . .
10 ADD https://cdn.zoom.us/prod/5.10.0.2450/zoom_amd64.deb .
11 EXPOSE 3000
12 USER node
13 CMD ["node", "index.js"]
```

Use multi-stage builds for small secure images

Use multi-stage builds in order to produce smaller and cleaner images, thus minimizing the attack surface for bundled docker image dependencies.



multi-stage-dockerfile x

multi-stage-dockerfile > ...

```
1 FROM golang AS builder
2 LABEL maintainer="Akshay Ithape"
3
4 RUN mkdir /app
5 WORKDIR /app
6
7 COPY hello.go .
8 RUN go build hello.go
9
10 FROM scratch
11 COPY --from=builder /app/hello .
12
13 ENTRYPOINT [ "./hello" ]
```

hello.go x

hello.go

```
1 package main
2
3 import "fmt"
4
5 func main () {
6     fmt.Println("Hello, world!")
7 }
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL

~/docker-test

> docker run test-hello:v1

Hello, world!

at 10:52:05 am

~/docker-test

> docker image ls test-hello:v1

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
test-hello	v1	77b8b8d90863	37 seconds ago	1.75MB

at 10:52:11 am

~/docker-test

>

at 10:52:13 am

zsh

zsh

Use a linter

Enforce Dockerfile best practices automatically by using a static code analysis tool such as **hadolint** linter, that will detect and alert for issues found in a Dockerfile.

```
DL4000 Specify a maintainer of the Dockerfile
DL3006 Always tag the version of an image explicitly.

1 FROM debian

SC1007 Remove space after = if trying to assign a value (for empty string, use var='' ... ).
SC2154 node_version is referenced but not assigned.
DL3009 Delete the apt-get lists after installing something

2 RUN node_version= "0.10" \
3   && apt-get update && apt-get -y install nodejs="$node_version"
4 COPY package.json usr/src/app

DL3003 Use WORKDIR to switch to a directory

5 RUN cd /usr/src/app \
6   && npm install node-static
7

DL3011 Valid UNIX ports range from 0 to 65535

8 EXPOSE 80000
9 CMD ["npm", "start"]
```

Reference Link

- <https://snyk.io/blog/10-docker-image-security-best-practices/>
- <https://www.docker.com/blog/docker-and-snyk-extend-partnership-to-docker-official-and-certified-images/>
- <https://www.docker.com/blog/improve-the-security-of-hub-container-images-with-automatic-vulnerability-scans/>
- <https://www.docker.com/blog/advanced-dockerfiles-faster-builds-and-smaller-images-using-buildkit-and-multistage-builds/>
- <https://betterprogramming.pub/docker-content-trust-security-digital-signatures-eeae9348140d>
- <https://snyk.io/plans/>
- <https://aquasecurity.github.io/trivy/v0.18.3/installation/>

Thank You EveryOne

Be In Touch

Linkedin : <https://www.linkedin.com/in/akshayithape/>

Gmail : ithapeakshay.02@gmail.com

GitHub : <https://github.com/akshayithape-devops>

Medium : <https://akshayithape.medium.com/>